Responsible AI and academic integrity

- Attempt exercises first; use graduated hints and explain every submitted line.
- Credit AI assistance and source authors. Follow your institution's disclosure policy.
- Run generated code in a controlled environment; inspect imports, file writes and network calls first.
- Never use eval on learner input; never run unknown installation commands blindly.
- Keep passwords, access tokens and student personal information out of prompts and GitHub.
- Do not invent tool usage, screenshots, learner scores, citations or execution claims.
- Do not republish publisher text. Summarize in your own words and link to the source.
- AI feedback is fallible; retain human review and allow learners to challenge explanations.
- Use pseudonyms and consent for evaluation. Do not commit identifiable learner data.